

Topics Covered : Distributed Databases: Introduction to Distributed Database Systems, Distributed Database System Architecture;

Dr. V.C.V. Rao

Office @Pune University Campus, Pune, India

*CDAC is a R&D Institute of Dept. of Electronics & Information Technology (DeitY),
Ministry of Communications & Information Technology (MCIT); Govt. of India*

NIT-Warangal:

Dept. Of Comp. Sc & Engg.

July – Nov 2023

Course Name (CS5305) : Advanced DataBases

Course Name (CS5309) : Advanced DataBases

Advanced Data Bases (References & Source for Presentation)

The Presentation is prepared based on Author's experience on various research projects on Distributed Computing – Distributed Data Bases (NoSQL) well as references given in this presentation.

Source :

Text Books, Research Articles, Web Sites as indicated in many slides and References of this presentation

1. M T Ozsú, Patrick Valduriez, Principles of Distributed Database Systems, Prentice Hall, 3rd Edition 2011

Syllabus:

UNIT 1 : Distributed Databases: Introduction to Distributed Database Systems, Distributed Database System Architecture;

UNIT 2 : Top-Down Approach, Distributed Database Design Issues, Fragmentation, Allocation, Database Integration, Bottom-up approach, Schema Matching, Schema Integration, Schema Mapping;

UNIT 3: Data and Access Control, View Management, Data Security; Query processing problem, Objectives of Query processing, Complexity of Relational Algebra Operations, Characterization of Query Processors, Layers of Query Processing;

UNIT 4 : Query Decomposition, Normalization, Analysis, Elimination of Redundancy and Rewriting;

UNIT 5 : Localization of Distributed Data, Reduction for primary Horizontal, Vertical, derived Fragmentation;

UNIT 6 : Distributed Query Execution, Query Optimization, Join Ordering, Static& Dynamic Approach, Semi-joins, Hybrid Approach;

Syllabus:

M.C.A 2nd Year, July 2023

UNIT 7 : Taxonomy of Concurrency control Mechanisms, Lock-Based Concurrency Control, Timestamp-Based Concurrency Control, Optimistic Concurrency Control, Deadlock Management;

UNIT 8 : Heterogeneity issues Advanced Transaction Models, Distributed systems 2PC& 3PC protocols, Replication protocols, Replication and Failures, HotSpares;

UNIT 9: Parallel Databases: Introduction to Parallel Databases, Parallel Database System Architectures, Parallel Data Placement, Full Partitioning; Parallel Query Processing, Query Parallelism;

UNIT 10 : Parallel Query Optimization, Search Space, Cost Model, Search Strategy; Load Balancing.

Text Books/ Reference Books / Online Resources:

M T Ozsu, Patrick Valduriez, Principles of Distributed Database Systems, Prentice Hall, 1999.

S. Ceri and G. Pelagatti, Distributed Database System Principles and Systems, MGH, 1985.

Course Name : Advanced databases

Topics Covered in this Lecture

UNIT 1 : Distributed Databases: Introduction to Distributed Database Systems, Distributed Database System Architecture;

UNIT 2 : Top-Down Approach, Distributed Database Design Issues, Fragmentation, Allocation, Database Integration, Bottom-up approach, Schema Matching, Schema Integration, Schema Mapping;

Chapter-02 & Chapter-03 of Reference Book No. 1

Background

1. Two technological bases for distributed database technology: **database management** and **computer networks**
2. Overview of the concepts in these two fields that **are more important from** the perspective of distributed database technology.

Overview of Relational DBMS

1. Define the terminology and framework used
 2. Most of the distributed database technology has been developed using the relational mode
- ❖ Relational Database Concepts
 - Normalisation
 - Relational Data Languages
 - ❖ Review of Computer Networks

Relational Database Concepts

A database is a structured collection of data related to some real-life phenomena that we are trying to model. A relational database is one where the database structure is in the form of tables. Formally, a relation **R** defined over **n** sets $D_1, D_2, :::, D_n$ (not necessarily distinct) is a set of **n**-tuples (or simply tuples)

$$\langle d_1, d_2, :::, d_n \rangle$$

such that

$$d_1 \in D_1, d_2 \in D_2, :, d_n \in D_n.$$

Relational Database Concepts

Example : Use a database that models an engineering company. The entities to be modelled are the employees (**EMP**) and projects (**PROJ**).

For each employee, keep track of the

- employee number (**ENO**),
- name (**ENAME**),
- title in the company (**TITLE**),
- Salary (**SAL**),
- identification number of the project(s) the employee is working on (**PNO**),
- responsibility within the project (**RESP**), and duration of the assignment to the project (**DUR**) in months.
- Similarly, for each project, to store the project number (**PNO**), the project name (**PNAME**), & the project budget (**BUDGET**).

Relational Database Concepts

Sample Database Scheme

EMP

ENO	ENAME	TITLE	SAL	PNO	RESP	DUR
-----	-------	-------	-----	-----	------	-----

PROJ

PNO	PNAME	BUDGET
-----	-------	--------

The *relation schemas* for this database can be defined as follows:

EMP(ENO, ENAME, TITLE, SAL, PNO, RESP, DUR)

PROJ(PNO, PNAME, BUDGET)

In relation scheme **EMP**, there are seven attributes: **ENO, ENAME, TITLE, SAL, PNO, RESP, DUR**. The values of **ENO** come from the domain of all valid employee numbers, say D_1 , the values of **ENAME** come from the domain of all valid names, say D_2 , and so on. Note that each attribute of each relation does not have to come from a distinct domain.

Relational Database Concepts

- Various attributes within a relation or from a number of relations may be defined over the same domain
- The **key** of a relation scheme is the minimum non-empty subset of its attributes such that the values of the attributes comprising the **key** uniquely identify each tuple of the relation.
- The attributes that make up **key** are called prime attributes. The superset of a **key** is usually called a **superkey**. Thus in our example the **key** of **PROJ** is **PNO**, and that of **EMP** is the set (**ENO, PNO**).

Relational Database Concepts

- Each relation has at least one **key**. Sometimes, there may be more than **one possibility** for the key.
- In such cases, each alternative is considered **a candidate key**, and one of the candidate keys is chosen as the **primary key**, which we denote by **underlining**.
- The number of attributes of a relation defines its **degree**, whereas the number of tuples of the relation defines its **cardinality**.
- In tabular form, the example database consists of two tables, as shown in Figure. The columns of the tables correspond to the attributes of the relations; if there were any information entered as the rows, they would correspond to the tuples

Relational Database Concepts

- ❖ In tabular form, the example database consists of two tables, as shown in Figure.
 - The columns of the tables correspond to the attributes of the relations;
- ❖ if there were any information entered as the rows, they would correspond to the tuples.
- ❖ The empty table, showing the structure of the table, corresponds to the relation schema; when the table is filled with rows, it corresponds to a relation instance.

Relational Database Concepts

EMP

ENO	ENAME	TITLE	SAL	PNO	RESP	DUR
E1	J. Doe	Elect. Eng.	40000	P1	Manager	12
E2	M. Smith	Analyst	34000	P1	Analyst	24
E2	M. Smith	Analyst	34000	P2	Analyst	6
E3	A. Lee	Mech. Eng.	27000	P3	Consultant	10
E3	A. Lee	Mech. Eng.	27000	P4	Engineer	48
E4	J. Miller	Programmer	24000	P2	Programmer	18
E5	B. Casey	Syst. Anal.	34000	P2	Manager	24
E6	L. Chu	Elect. Eng.	40000	P4	Manager	48
E7	R. Davis	Mech. Eng.	27000	P3	Engineer	36
E8	J. Jones	Syst. Anal.	34000	P3	Manager	40

PROJ

PNO	PNAME	BUDGET
P1	Instrumentation	150000
P2	Database Develop.	135000
P3	CAD/CAM	250000
P4	Maintenance	310000

Figure. Sample Database Instance

Overview of Relational DBMS - Normalization

The aim of normalization is to eliminate various anomalies (or undesirable aspects) of a relation in order to obtain “better” relations. The following four problems might exist in a relation scheme:

- Repetition anomaly
- Update anomaly
- Insertion anomaly
- Deletion anomaly

Normalization transforms arbitrary relation schemes into ones without these problems. A relation with one or more of the above mentioned anomalies is split into two or more relations of a higher normal form.

Overview of Relational DBMS-Relational Data Languages

Data manipulation languages developed for the relational model (commonly called ***query languages***) fall into two fundamental groups: ***relational algebra languages*** and ***relational calculus languages***. The difference between them is based on how the user query is formulated.

- The relational algebra is procedural in that the user is expected to specify, using certain high-level operators, how the result is to be obtained.
- The relational calculus, on the other hand, is non-procedural; the user only specifies the relationships that should hold in the result.

Relational Data Languages - Relational Algebra

Relational Algebra

- Relational algebra consists of a set of operators that operate on relations.
- Each operator takes one or two relations as operands and produces a result relation, which, in turn, may be an operand to another operator.
- These operations permit the querying and updating of a relational database.

Relational Data Languages - Relational Algebra

Normalized Relations

EMP

ENO	ENAME	TITLE
E1	J. Doe	Elect. Eng
E2	M. Smith	Syst. Anal.
E3	A. Lee	Mech. Eng.
E4	J. Miller	Programmer
E5	B. Casey	Syst. Anal.
E6	L. Chu	Elect. Eng.
E7	R. Davis	Mech. Eng.
E8	J. Jones	Syst. Anal.

ASG

ENO	PNO	RESP	DUR
E1	P1	Manager	12
E2	P1	Analyst	24
E2	P2	Analyst	6
E3	P3	Consultant	10
E3	P4	Engineer	48
E4	P2	Programmer	18
E5	P2	Manager	24
E6	P4	Manager	48
E7	P3	Engineer	36
E8	P3	Manager	40

PROJ

PNO	PNAME	BUDGET
P1	Instrumentation	150000
P2	Database Develop.	135000
P3	CAD/CAM	250000
P4	Maintenance	310000

PAY

TITLE	SAL
Elect. Eng.	40000
Syst. Anal.	34000
Mech. Eng.	27000
Programmer	24000

Relational Data Languages - Relational Algebra

Normalized Relations

There are five fundamental relational algebra operators and five others that can be defined in terms of these. The fundamental operators are **selection, projection, union, set difference, and Cartesian product.**

The first two of these operators are unary operators, and the last three are binary operators.

The additional operators that can be defined in terms of these fundamental operators are

intersection, θ - join, natural join, semi-join and division.

Relational Data Languages - Relational Algebra

Unary Relational Operations

SELECT (symbol: σ)

PROJECT (symbol: π)

RENAME (symbol: ρ)

Relational Algebra Operations From Set Theory

UNION (\cup)

INTERSECTION (\cap),

DIFFERENCE ($-$)

CARTESIAN PRODUCT (\times)

Binary Relational Operations

JOIN

DIVISION

Relational Calculus

In Relational calculus-based languages, instead of specifying how to obtain the result, one specifies what the result is by stating the relationship that is supposed to hold for the result.

Relational calculus languages fall into two groups: tuple relational calculus and domain relational calculus. The difference between the two is in terms of the primitive variable used in specifying the queries.

Chapter-03

Part-01

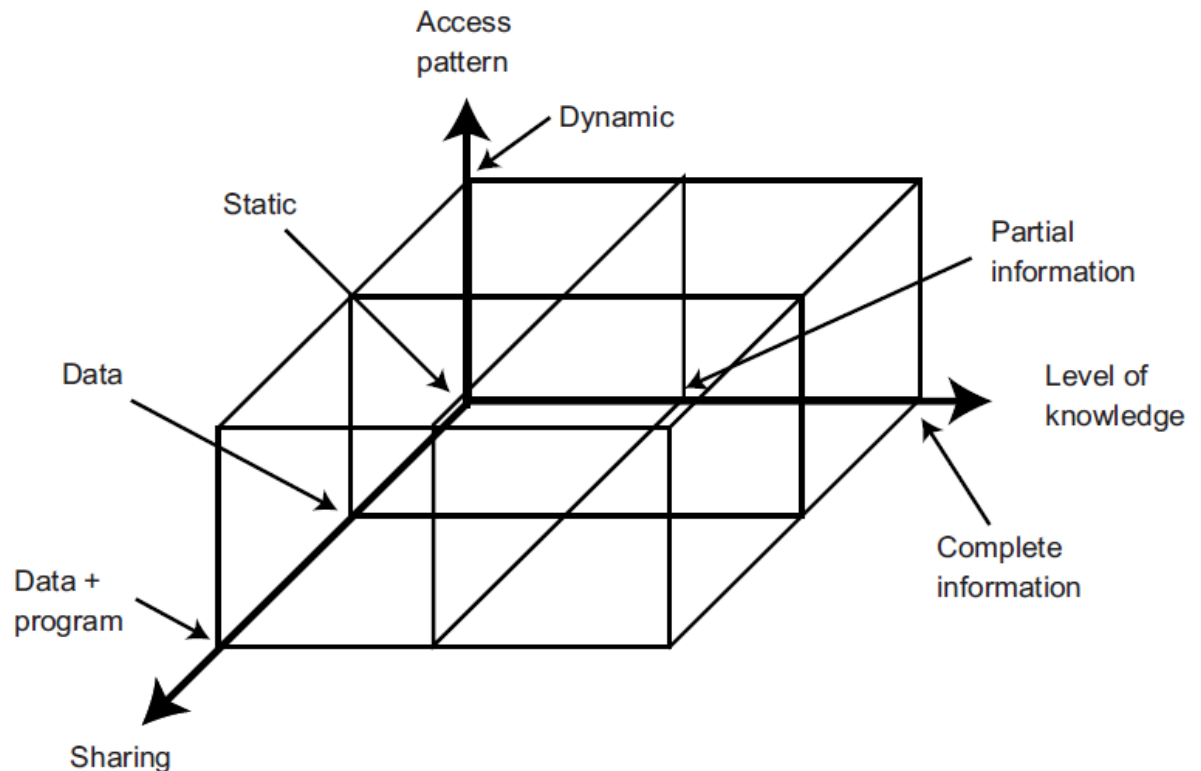
Distributed Database Design

1. Top-Down Design Process
2. Distribution Design Issues
3. Fragmentation
4. Allocation
5. Data base Integration

Distributed Database Design

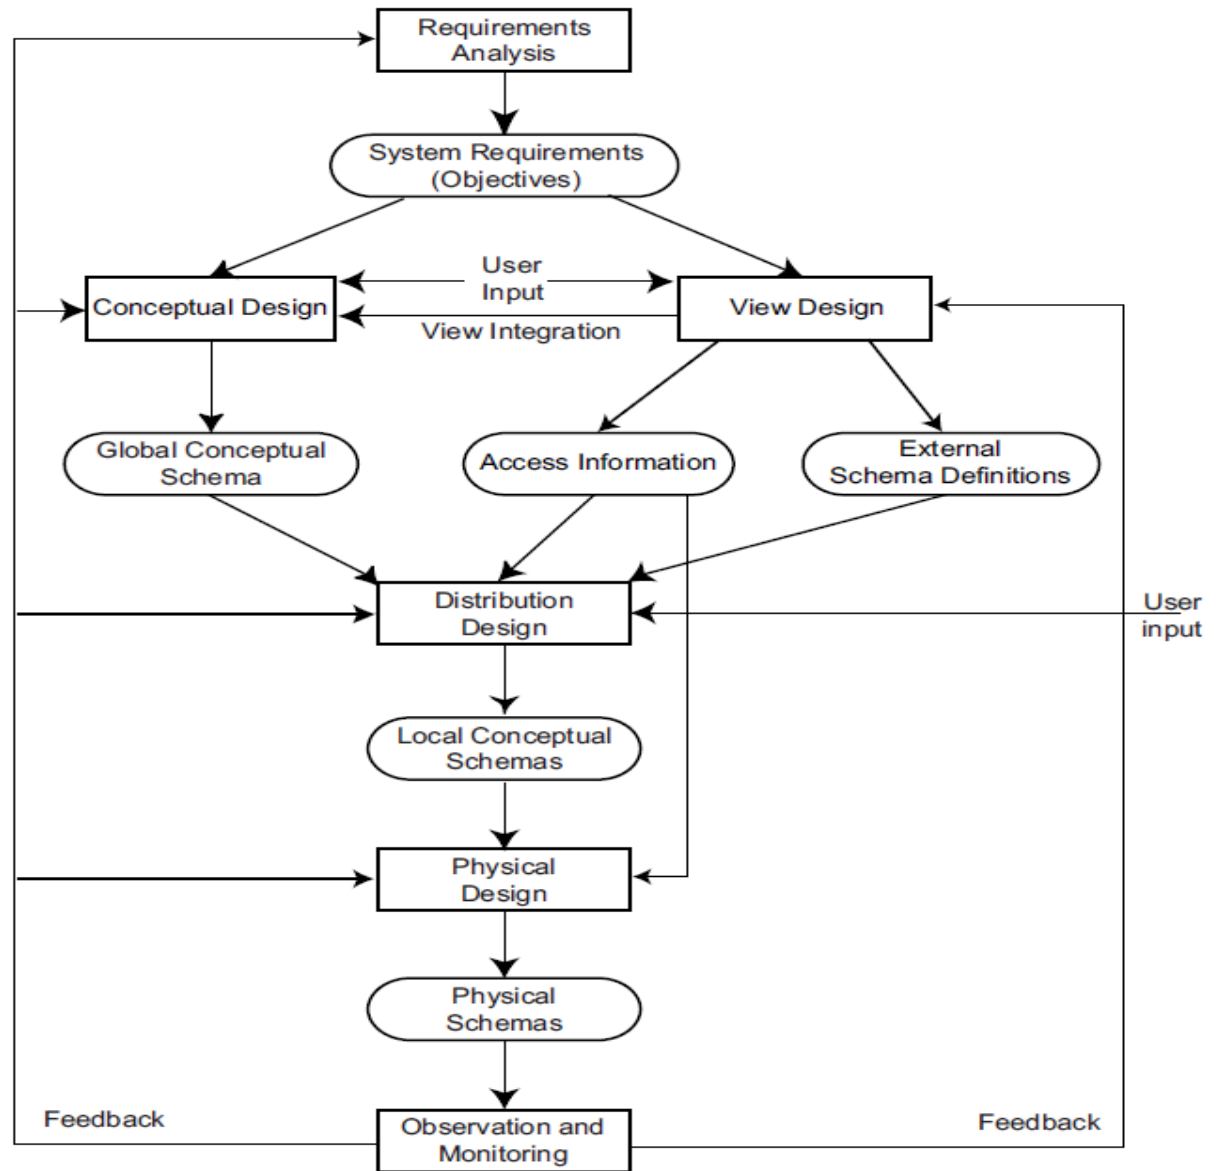
The organization of distributed systems can be investigated along three orthogonal dimensions (Figure)

1. Level of sharing
2. Behavior of access patterns
3. Level of knowledge on access pattern behaviour



Top Down Design Process

The organization of distributed systems can be investigated along three orthogonal dimensions (Figure)



Top Down Design Process

1. The conceptual design, on the other hand, is the process by which the enterprise is examined to determine entity types and relationships among these entities.
2. One can possibly divide this process into **two related activity groups**
 - **entity analysis and functional analysis** (This is all part of **Centralized Data Base Design**)

Top Down Design Process

1. The **global conceptual schema (GCS)** and access pattern information collected as a result of view design are inputs to the distribution design step.
2. The objective at this stage is to design the **local conceptual schemas (LCSs)** by distributing the entities over the sites of the distributed system
3. Rather than distributing relations, it is quite common to divide them into subrelations, called **fragments**, which are then distributed. Thus, the distribution design activity consists of two steps: **fragmentation** and **allocation**.

Distributed Design Issues

Questions should be addresses

1. Why fragment at all?
2. How should we fragment?
3. How much should we fragment?
4. Is there any way to test the correctness of decomposition?
5. How should we allocate?
6. What is the necessary information for fragmentation and allocation?

Reasons for fragmentation

Questions should be addresses

With respect to fragmentation, the important issue is the appropriate unit of distribution. A relation is not a suitable unit, for a number of reasons.

- **First**, application views are usually subsets of relations. (consider subsets of relations as distribution units.)
Second, if the applications that have views defined on a given relation reside at different sites, two alternatives can be followed (Replication at one site or all sites)
- **Finally**, the decomposition of a relation into fragments, each being treated as a unit, permits a number of transactions to execute concurrently

Reasons for fragmentation

Fragmentation raises difficulties (Problem 1 & Problem 2)

Problem 1 :

- If the applications have conflicting requirements that prevent decomposition of the relation into mutually exclusive fragments
- Performance Degradation - Example necessary to retrieve data from two fragments and then take their join, which is costly
- Minimizing distributed joins is a fundamental fragmentation issue.

Reasons for fragmentation

Problem 2 :

- The second problem is related to semantic data control, specifically to integrity checking.
- As a result of fragmentation, attributes participating in a dependency may be decomposed into different fragments that might be allocated to different sites.

Fragmentation Alternatives

Problem 2 :

- Relation instances are essentially tables, so the issue is one of finding alternative ways of dividing a table into smaller ones. There are clearly two alternatives for this: dividing it horizontally or dividing it vertically.

we use a modified version of the relational database scheme developed .

We have added to the **PROJ** relation a new attribute (**LOC**) that indicates the place of each project.

Fragmentation Alternatives

Example 1 : We have added to the **PROJ** relation a new attribute (**LOC**) that indicates the place of each project.

- ❖ Figure (A) depicts the database instance we will use. Figure (B) shows the **PROJ** relation of Figure (A) divided horizontally into two relations. Sub-relation
 - **PROJ1** contains information about projects whose budgets are less than \$200,000, whereas
 - **PROJ2** stores information about projects with larger budgets.

Fragmentation Alternatives

Example 1 : We have added to the **PROJ** relation a new attribute (**LOC**) that indicates the place of each project.

❖ Figure (A) depicts the database instance we will use.
Figure (B) shows the **PROJ** relation of Figure (A) divided horizontally into two relations. Sub-relation

- **PROJ1** contains information about projects whose budgets are less than \$200,000, whereas
- **PROJ2** stores information about projects with larger budgets.

Note that in the case of horizontal fragmentation, the “item” typically refers to a tuple, while in the case of vertical fragmentation, it refers to an attribute.

Fragmentation Alternatives

Example 1 : We have added to the **PROJ** relation a new attribute (**LOC**) that indicates the place of each project.

PROJ₁

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal
P2	Database Develop.	135000	New York

PROJ₂

PNO	PNAME	BUDGET	LOC
P3	CAD/CAM	255000	New York
P4	Maintenance	310000	Paris

Fig. (B) Example of Horizontal Partitioning

Fragmentation Alternatives

Example 3.2. Figure (C) shows the PROJ relation of Figure (A) partitioned vertically into two subrelations **PROJ1** and **PROJ2**.

- **PROJ1** contains only the information about project budgets, whereas
- **PROJ2** contains project names and locations. It is important

Note that in the case of horizontal fragmentation, the “item” typically refers to a tuple, while in the case of vertical fragmentation, it refers to an attribute.

Fragmentation Alternatives

Example 3.2. Figure (C) shows the PROJ relation of Figure (A) partitioned vertically into two subrelations **PROJ1** and **PROJ2**.

PROJ ₁		PROJ ₂		
PNO	BUDGET	PNO	PNAME	LOC
P1	150000	P1	Instrumentation	Montreal
P2	135000	P2	Database Develop.	New York
P3	250000	P3	CAD/CAM	New York
P4	310000	P4	Maintenance	Paris

Figure C. Example of Vertical Partitioning

In general, the applications need to be characterized with respect to a number of parameters

Fragmentation Alternatives

❖ **Degree of Fragmentation** (Performance of query execution)

❖ **Correctness Rules of Fragmentation (**

- Completeness,
- Reconstruction, &
- Disjoint Ness

Allocation Alternatives

- ❖ Assuming that the database is fragmented properly, one has to decide on the allocation of the fragments to various sites on the network.
- ❖ When data are allocated, it may either be replicated or maintained as a single copy.
 - The reasons for replication are reliability and efficiency of read-only queries.
 - The decision regarding replication is a trade-off that depends on the ratio of the read-only queries to the update queries. This decision affects almost all of the distributed DBMS algorithms and control functions.

Allocation Alternatives

- ❖ A non-replicated database (commonly called a **partitioned database**) contains fragments that are allocated to sites, and there is only one copy of any fragment on the network.
- ❖ In case of replication, either the database exists in its entirety at each site (**fully replicated database**), or fragments are distributed to the sites in such a way that copies of a fragment may reside in multiple sites (**partially replicated database**).

Note : In the latter the number of copies of a fragment may be an input to the allocation algorithm or a decision variable whose value is determined by the algorithm.

Allocation Alternatives

Comparison of Replication Alternatives

	Full replication	Partial replication	Partitioning
QUERY PROCESSING	Easy	← Same difficulty →	
DIRECTORY MANAGEMENT	Easy or nonexistent	← Same difficulty →	
CONCURRENCY CONTROL	Moderate	Difficult	Easy
RELIABILITY	Very high	High	Low
REALITY	Possible application	Realistic	Possible application

Fragmentation

❖ Fragmentation Strategies & Algorithms

Horizontal Fragmentation

Partitions a relation along its tuples.

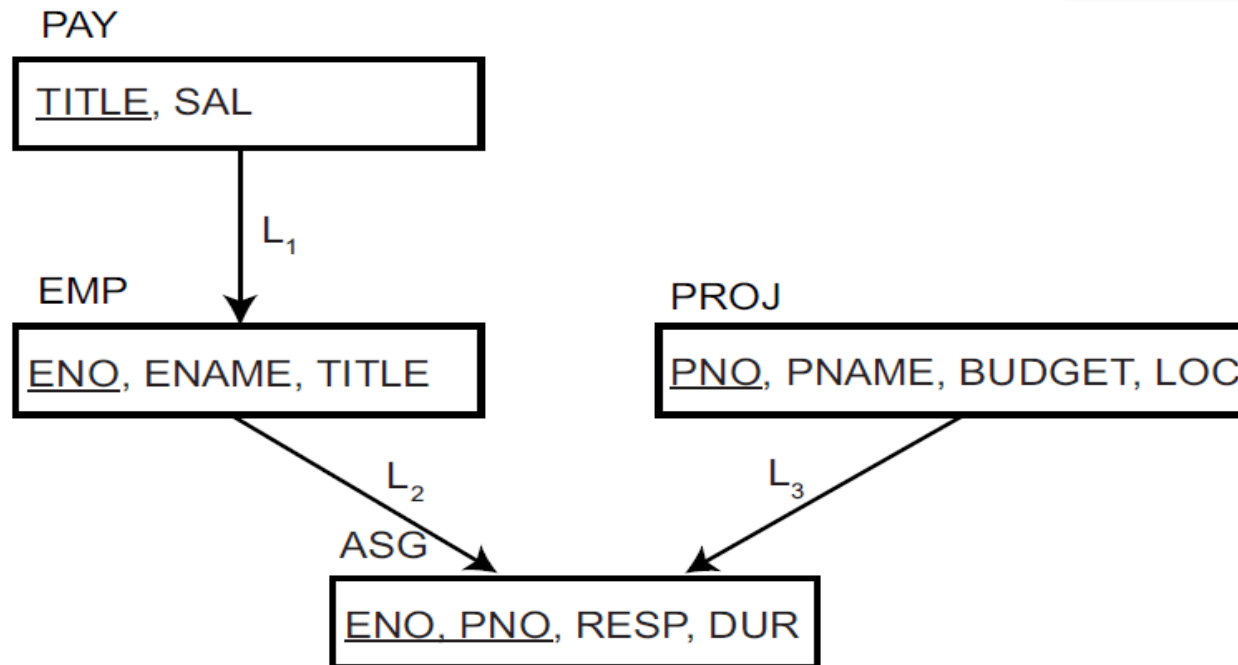
- Primary horizontal fragmentation &
- Derived horizontal fragmentation

❖ Information Requirements of Horizontal Fragmentation

- Database Information
- Application Information.

Fragmentation

❖ Fragmentation Strategies & Algorithms



Expression of Relationships Among Relations Using Links

Primary Horizontal Fragmentation

A **primary horizontal fragmentation** is defined by a selection operation on the owner relations of a database schema.

Therefore, given relation R, its horizontal fragments are given by

$$R_i = \sigma_{F_i}(R); 1 \leq i \leq w$$

Where F_i is the selection formula used to obtain fragment R_i (also called the fragmentation predicate). **Algorithms are required**

The decomposition of relation PROJ into horizontal fragments PROJ₁ and PROJ₂ in Example (Earlier) is defined as follows:

$$PROJ_1 = \sigma_{BUDGET \leq 200000}(PROJ)$$

$$PROJ_2 = \sigma_{BUDGET > 200000}(PROJ)$$

Primary Horizontal Fragmentation

PROJ₁

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal

PROJ₂

PNO	PNAME	BUDGET	LOC
P2	Database Develop.	135000	New York
P3	CAD/CAM	250000	New York

PROJ₃

PNO	PNAME	BUDGET	LOC
P4	Maintenance	310000	Paris

Primary Horizontal Fragmentation of Relation PROJ

$$Pr = \{ \text{LOC} = \text{"Montreal"}, \text{LOC} = \text{"New York"}, \text{LOC} = \text{"Paris"}, \\ \text{BUDGET} \leq 200000, \text{BUDGET} > 200000 \}$$

Primary Horizontal Fragmentation

An iterative algorithm that would generate a **complete and minimal set of predicates** P'_r given a set of simple predicates P_r .

PAY₁

TITLE	SAL
Mech. Eng.	27000
Programmer	24000

PAY₂

TITLE	SAL
Elect. Eng.	40000
Syst. Anal.	34000

Horizontal Fragmentation of Relation PAY

Primary Horizontal Fragmentation

Horizontal Partitioning of Relation PROJ

PROJ₁

PNO	PNAME	BUDGET	LOC
P1	Instrumentation	150000	Montreal

PROJ₃

PNO	PNAME	BUDGET	LOC
P2	Database Develop.	135000	New York

PROJ₄

PNO	PNAME	BUDGET	LOC
P3	CAD/CAM	250000	New York

PROJ₆

PNO	PNAME	BUDGET	LOC
P4	Maintenance	310000	Paris

Derived Horizontal Fragmentation

A derived horizontal fragmentation is defined on a member relation of a link according to a selection operation specified on its owner.

It is important to remember two points.

- ❖ First, the link between the owner and the member relations is defined as an **equi-join**.
- ❖ Second, an **equi-join** can be implemented by means of **semijoins**.

Checking for Correctness

- ❖ **Completeness** : The completeness of a primary horizontal fragmentation is based on the selection predicates used.
- ❖ **Reconstruction** : Reconstruction of a global relation from its fragments is performed by the union operator in both the primary and the derived horizontal fragmentation.

Thus, for a relation R with fragmentation

$$F_R = f\{R_1; R_2; : : : ; R_w\},$$

$$R = \bigcup R_i, \quad \forall R_i \in F_R$$

- ❖ **Disjointness.**

Derived Horizontal Fragmentation

EMP₁

ENO	ENAME	TITLE
E3	A. Lee	Mech. Eng.
E4	J. Miller	Programmer
E7	R. Davis	Mech. Eng.

EMP₂

ENO	ENAME	TITLE
E1	J. Doe	Elect. Eng.
E2	M. Smith	Syst. Anal.
E5	B. Casey	Syst. Anal.
E6	L. Chu	Elect. Eng.
E8	J. Jones	Syst. Anal.

Derived Horizontal Fragmentation of Relation EMP

Vertical Fragmentation

- ❖ The objective of vertical fragmentation is to partition a relation into a set of smaller relations so that many of the user applications will run on only one fragment.
- ❖ In this context, an “optimal” fragmentation is one that produces a fragmentation scheme which minimizes the execution time of user applications that run on these fragments.
- ❖ a vertical fragmentation of a relation **R** produces fragments **$R_1, R_2, : : : ; R_r$** , each of which contains a subset of R's attributes as well as the primary key of **R**.

Vertical Fragmentation

- ❖ Vertical fragmentation has been investigated within the context of centralized database systems as well as distributed ones.
- ❖ Its motivation within the centralized context is as a design tool, which allows the user queries to deal with smaller relations, thus causing a smaller number of page accesses
- ❖ Vertical partitioning is inherently more complicated than horizontal partitioning. This is due to the total number of alternatives that are available.

Vertical Fragmentation

- ❖ Two types of heuristic approaches exist for the vertical fragmentation of global relations:

Grouping and Splitting

- ❖ Information Requirements of Vertical Fragmentation (Based on Application)

Vertical Fragmentation

- ❖ Consider relation PROJ of Figure (Earlier) Assume that the following applications are defined to run on this relation. In each case we also give the SQL specification.
- ❖ **q1:** Find the budget of a project, given its identification number.

```
SELECT BUDGET  
FROM PROJ  
WHERE PNO=Value
```


Vertical Fragmentation

- ❖ **q2:** Find the names and budgets of all projects.

```
SELECT PNAME, BUDGET  
FROM   PROJ
```

- ❖ **q3:** Find the names of projects located at a given city.

```
SELECT PNAME  
FROM   PROJ  
WHERE  LOC=Value
```

- ❖ **q4:** Find the total project budgets for each city.

```
SELECT SUM (BUDGET)  
FROM   PROJ  
WHERE  LOC=Value
```

Vertical Fragmentation

- ❖ According to these four applications, the attribute usage values can be defined. As a notational convenience, we let

$A_1 = \text{PNO}$, $A_2 = \text{PNAME}$, $A_3 = \text{BUDGET}$, and $A_4 = \text{LOC}$.

- ❖ The usage values are defined in matrix form (Figure), where entry (i, j) denotes $\text{use}(\mathbf{q}_i, \mathbf{A}_j)$.
- ❖ Attribute usage values are not sufficiently general to form the basis of attribute splitting and fragmentation.

Vertical Fragmentation

❖ Attribute Affinity matrix (AA)

	A_1	A_2	A_3	A_4
q_1	1	0	1	0
q_2	0	1	1	0
q_3	0	1	0	1
q_4	0	0	1	1

Vertical Fragmentation

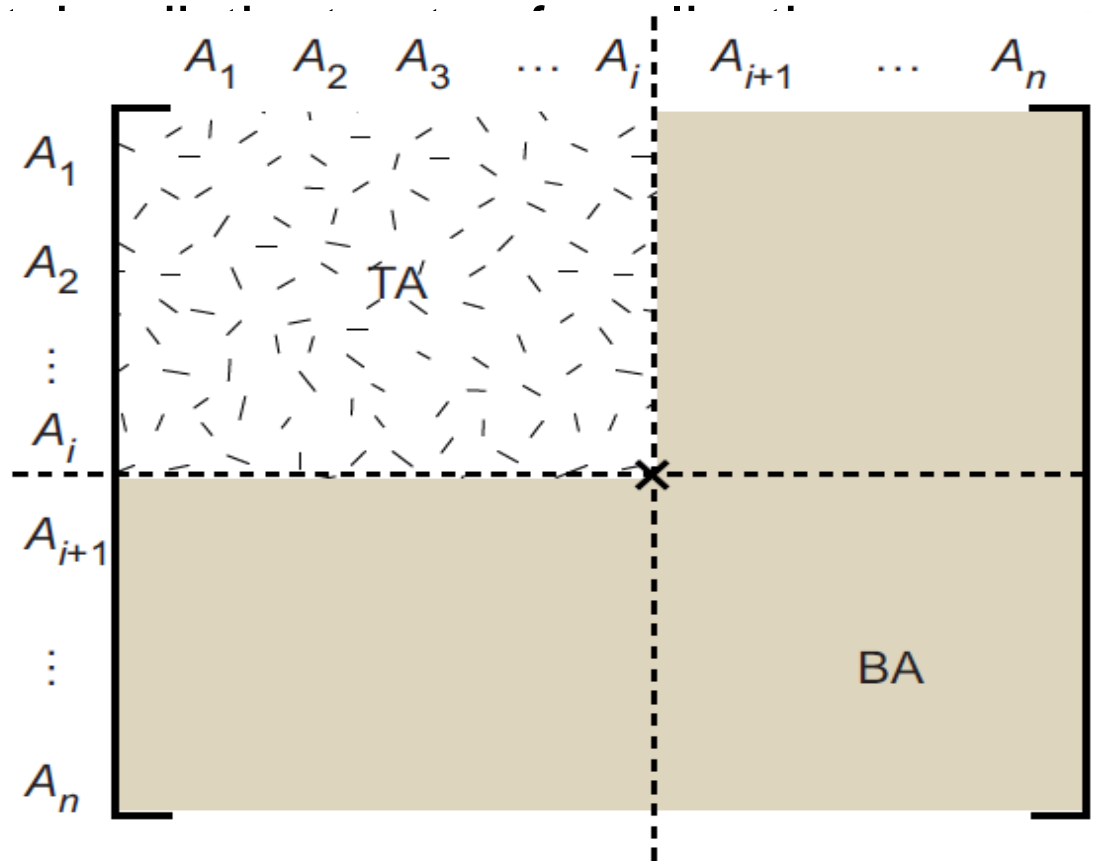
- ❖ **Clustering Algorithm** : The fundamental task in designing a vertical fragmentation algorithm is to find some means of grouping the attributes of a relation based on the attribute affinity values in AA. (attribute affinity matrix (AA).)
- 1. It is designed specifically to determine groups of similar items as opposed to, say, a linear ordering of the items (i.e., it clusters the attributes with larger affinity values together, and the ones with smaller values together).
- 2. The final groupings are insensitive to the order in which items are presented to the algorithm.
- 3. The computation time of the algorithm is reasonable: $O(n^2)$, where n is the number of attributes.
- 4. Secondary interrelationships between clustered attribute groups are identifiable

Vertical Fragmentation

- ❖ **Partitioning Algorithm** : The objective of the splitting activity is to find sets of attributes that are accessed solely, or for the most part, by distinct sets of applications
- ❖ For example, if it is possible to identify two attributes, A_1 and A_2 , which are accessed only by application q_1 , and attributes A_3 and A_4 , which are accessed by, say, two applications q_2 and q_3 , it would be quite straightforward to decide on the fragments.
- ❖ The task lies in finding an algorithmic method of identifying these groups

Vertical Fragmentation

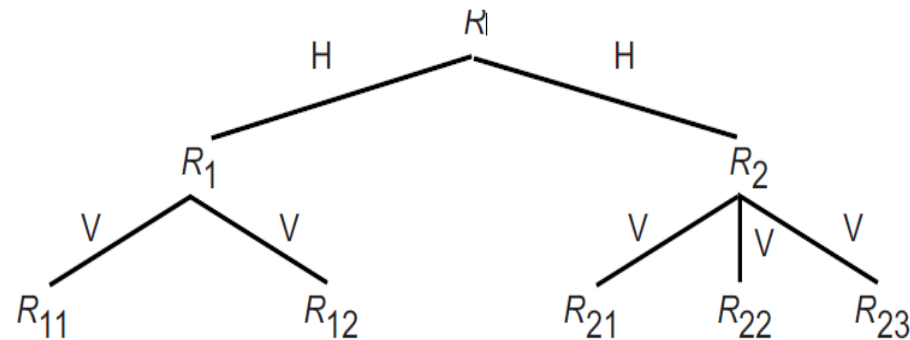
- ❖ **Partitioning Algorithm** : The objective of the splitting activity is to find sets of attributes that are accessed solely, or for the most part



Locating a Splitting Point

Hybrid Fragmentation

- ❖ In most cases a simple horizontal or vertical fragmentation of a database schema will not be sufficient to satisfy the requirements of user applications.
- ❖ In this case a vertical fragmentation may be followed by a horizontal one, or vice versa, producing a treestructured partitioning



Hybrid Fragmentation

Allocation

- ❖ Assume that there are a set of fragments

$$\mathbf{F} = \{\mathbf{F}_1, \mathbf{F}_2, : : : , \mathbf{F}_n\}$$

and a distributed system consisting of sites

$$\mathbf{S} = \{\mathbf{S}_1, \mathbf{S}_2, : : : , \mathbf{S}_m\}$$
 on which a set of applications

$$\mathbf{Q} = \{\mathbf{q}_1, \mathbf{q}_2, : : : , \mathbf{q}_q\}$$
 is running.

- ❖ The allocation problem involves finding the “optimal” distribution of \mathbf{F} to \mathbf{S} .
- ❖ The optimality can be defined with respect to two measures :
Minimal Cost and **Performance**.

Allocation

- ❖ Information Requirements
- ❖ Application Information
- ❖ Site Information
- ❖ Allocation Model
- ❖ Constraints
- ❖ Total Cost
- ❖ Solution Methods

Data Directory

- ❖ The distributed database schema needs to be stored and maintained by the system.
- ❖ This information is necessary during distributed query optimization,
- ❖ The schema information is stored in a data dictionary/directory, also called a catalog or simply a directory. A directory is a meta-database that stores a number of information

Chapter-03

Part-02